



# **Lying with Truths: Open-Channel Multi-Agent Collusion for Belief Manipulation via Generative Montage**

**Jinwei Hu<sup>1</sup>, Xinmiao Huang<sup>1</sup>, Youcheng Sun<sup>2</sup>, Yi Dong<sup>1\*</sup>, Xiaowei Huang<sup>1</sup>**

<sup>1</sup>University of Liverpool, United Kingdom

<sup>2</sup>Mohamed bin Zayed University of Artificial Intelligence, UAE

{jinwei.hu, xinmiao.huang, yi.dong, xiaowei.huang}@liverpool.ac.uk

youcheng.sun@mbzuai.ac.ae

HUMANE Lab, Taegyun Kim, M.S. student

ACL 2026

2026.09.18

# LLMs and narrative coherence


---

- LLMs seek coherent explanations from fragmented or ambiguous inputs
- They infer causal links between unrelated facts and form false beliefs
- Those beliefs can make agents unwitting amplifiers of misinformation

# Prior collusion: covert channel


---

- Existing research focuses on hidden communication paths between colluding agents
  - Attackers: use covert channels or backdoors to deliver malicious content
  - Detectors: seek to intercept the concealed message before it reaches the victim



# Can truths make an LLM infer a false story?

- Attackers distribute fragmented, truthful evidence through public channel
- A victim may connect unrelated facts into a spurious causal story




# Problem formulation


---

- Cognitive collusion attacks: construct an optimal evidence stream that maximizes the victim's posterior belief without fabricating
  - Local truth (LT): fully consistent with the ground truth
  - Global lie (GL): successfully steers induces stronger belief in a fabricated hypothesis
- Colluder
  - Explicit colluders: intentionally optimize deceptive objectives
  - Implicit colluders: unintentionally amplify deception

# Generative montage framework



# Generative montage framework




# Adversarial narrative production

---

- Writer: select evidence and construct narrative
- Editor: strategically arranging fragments with subtle semantic associations and temporal proximities
- Director: simulate the victim's belief update and evaluate deceptive effectiveness while preserving factual integrity
  - Decision: accept / reject / revise

# Generative montage framework




# Decentralized injection via publisher

---

- Decompose the approved montage sequence into individually fragments
- Assign fragments to coordinate Sybil publishers across public channel
- Trigger belief updating while reducing the appearance of coordinated manipulation

# Generative montage framework




# Cognitive steering via overthinking

---

- Exploit the victim's tendency to resolve ambiguity through a coherent causal narrative
- Victim internalizes the coerced inference as a self-derived conclusion

# Generative montage framework



# Cascade misinformation

---

- Downstream agent updates its belief from victim outputs
- Repeated victim endorsements amplify confidence in fabricated hypothesis
- Adversarial narrative spreads as an validated consensus

# Example

---

| Event                                                                                                        | Writer Output                                                                                                     | Editor Output                                                                                        | Victim Response                                      | Conf.       | Down.                                |
|--------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|------------------------------------------------------|-------------|--------------------------------------|
| <b>Charlie Hebdo</b><br><i>H<sub>f</sub>: Ahmed Merabet was the first victim of the Charlie Hebdo attack</i> | Officer Merabet stationed outside...confrontation occurred...#JeSuisAhmed trended...death confirmed in reports... | <i>Temporal proximity:</i><br>Attack starts → Officer present → Hashtag emerges → Casualty (7 posts) | <b>TRUE</b> — "Timeline suggests early encounter..." | <b>0.92</b> | MV: Y (5/5)<br>AI: Y<br><b>GT: N</b> |

# Experiments

---

- Dataset: CoPHEME
  - Adapted from PHEME to separate factual evidence from target conclusion

```
{
  "tweet_id": "552825739075022848",
  "text": "If all white people aren't responsible for #NAACPBombing than neither should all Muslims be responsible for #CharlieHebdo. #JeSuisCharlie",
  "timestamp": "2015-01-07T13:55:01+00:00",
  "veracity": "non-rumour",
  "is_rumour": false,
  "has_image": false,
  "has_video": false,
  "media_urls": [],
  "is_source": true,
  "parent_id": null,
  "retweet_count": 261,
  "favorite_count": 139,
  "user_id": "1769646126",
  "user_verified": false,
  "source_path": "/your_path/PHEME/all-rnr-annotated-threads/charliehebdo-all-rnr-threads/non-rumours/552825739075022848/source-tweets/552825739075022848.json"
},
```

# Experiments

---

- Simulation framework
  - Colluding group
  - LLM-based analyst (Victim)
  - Downstream decision layer
    - Majority vote
    - AI judge

# Experiments

---

- Evaluation metrics
  - Attack success rate
  - Average confidence
  - High-confidence ASR
  - Downstream deception rate

# Overall attack effectiveness

- Achieves over 70% overall ASR exhibiting high susceptibility

| Victim Model            | Charlie Hebdo |      |      | Sydney Siege |      |      | Ferguson |      |      | Ottawa Shoot. |      |      | Germanwings |      |      | Putin Missing |      |      | Overall ASR |
|-------------------------|---------------|------|------|--------------|------|------|----------|------|------|---------------|------|------|-------------|------|------|---------------|------|------|-------------|
|                         | A             | C    | H    | A            | C    | H    | A        | C    | H    | A             | C    | H    | A           | C    | H    | A             | C    | H    |             |
| Proprietary Models      |               |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| GPT-4o-mini             | 81.7          | 0.83 | 67.4 | 92.1         | 0.82 | 70.3 | 79.5     | 0.81 | 59.0 | 86.7          | 0.85 | 78.2 | 74.5        | 0.85 | 60.0 | 66.7          | 0.82 | 53.3 | 83.1        |
| GPT-4o                  | 79.4          | 0.85 | 66.9 | 88.5         | 0.83 | 69.1 | 67.0     | 0.83 | 53.0 | 85.5          | 0.85 | 74.5 | 56.4        | 0.89 | 52.7 | 63.3          | 0.75 | 26.7 | 77.4        |
| GPT-4.1-nano            | 81.7          | 0.81 | 66.9 | 94.5         | 0.80 | 67.1 | 79.1     | 0.79 | 53.1 | 94.5          | 0.81 | 74.8 | 74.5        | 0.83 | 61.8 | 70.0          | 0.75 | 13.3 | 85.5        |
| GPT-4.1-mini            | 79.9          | 0.85 | 68.4 | 80.6         | 0.82 | 50.9 | 68.0     | 0.81 | 42.5 | 77.6          | 0.87 | 67.9 | 63.6        | 0.91 | 63.6 | 20.0          | 0.81 | 16.7 | 72.7        |
| GPT-4.1                 | 77.1          | 0.88 | 73.7 | 64.8         | 0.88 | 60.6 | 60.0     | 0.87 | 57.5 | 77.0          | 0.89 | 74.5 | 54.5        | 0.94 | 54.5 | 16.7          | 0.90 | 16.7 | 65.9        |
| Claude-3-Haiku          | 94.8          | 0.80 | 69.0 | 98.8         | 0.79 | 72.0 | 83.5     | 0.76 | 50.0 | 98.8          | 0.79 | 66.9 | 68.5        | 0.82 | 61.1 | 86.7          | 0.71 | 26.7 | 91.5        |
| Claude-3.5-Haiku        | 76.9          | 0.81 | 47.9 | 76.9         | 0.79 | 45.0 | 77.1     | 0.79 | 46.9 | 80.5          | 0.77 | 38.4 | 61.8        | 0.82 | 38.2 | 74.1          | 0.71 | 18.5 | 76.7        |
| Claude-4.5-Haiku        | 52.6          | 0.74 | 20.6 | 34.5         | 0.70 | 10.3 | 32.0     | 0.71 | 8.0  | 50.9          | 0.74 | 14.6 | 63.6        | 0.77 | 34.5 | 16.7          | 0.61 | 16.7 | 42.4        |
| Proprietary Avg.        | 78.0          | 0.82 | 60.1 | 78.8         | 0.80 | 55.7 | 68.3     | 0.80 | 46.2 | 81.4          | 0.82 | 61.2 | 64.7        | 0.85 | 53.3 | 51.8          | 0.76 | 23.6 | 74.4        |
| Open-Weights Models     |               |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| Qwen2.5-3B-Inst         | 54.7          | 0.88 | 51.2 | 59.4         | 0.88 | 53.1 | 52.3     | 0.88 | 49.2 | 64.4          | 0.88 | 60.0 | 45.3        | 0.90 | 43.4 | 31.0          | 0.78 | 20.7 | 55.4        |
| Qwen2.5-7B-Inst         | 67.8          | 0.86 | 62.1 | 82.4         | 0.84 | 68.5 | 62.0     | 0.85 | 56.5 | 64.0          | 0.86 | 54.3 | 65.5        | 0.86 | 61.8 | 36.7          | 0.80 | 26.7 | 67.1        |
| Qwen2.5-14B-Inst        | 71.4          | 0.85 | 59.5 | 81.3         | 0.82 | 62.6 | 60.2     | 0.83 | 46.9 | 85.9          | 0.85 | 69.3 | 53.9        | 0.89 | 50.0 | 69.0          | 0.75 | 27.6 | 71.9        |
| DS-R1-Distill-Qwen-1.5B | 71.0          | 0.76 | 50.8 | 81.0         | 0.76 | 55.6 | 61.1     | 0.70 | 37.6 | 72.5          | 0.74 | 47.5 | 79.1        | 0.73 | 51.2 | 75.0          | 0.62 | 33.3 | 71.6        |
| DS-R1-Distill-Qwen-7B   | 74.9          | 0.88 | 66.3 | 92.1         | 0.85 | 75.2 | 76.0     | 0.87 | 69.5 | 81.2          | 0.89 | 74.5 | 66.0        | 0.92 | 60.4 | 66.7          | 0.86 | 60.0 | 79.2        |
| DS-R1-Distill-Qwen-14B  | 77.0          | 0.86 | 64.9 | 83.6         | 0.85 | 71.5 | 74.9     | 0.84 | 64.3 | 88.5          | 0.88 | 84.2 | 54.5        | 0.91 | 52.7 | 36.7          | 0.74 | 20.0 | 76.8        |
| Open-Weights Avg.       | 69.3          | 0.85 | 58.9 | 80.4         | 0.83 | 65.2 | 64.6     | 0.83 | 54.0 | 76.0          | 0.85 | 64.6 | 61.4        | 0.86 | 53.3 | 54.5          | 0.76 | 33.6 | 70.6        |

# Vulnerability of reasoning models

- Reasoning model exhibit higher vulnerability

| Victim Model            | Charlie Hebdo |      |      | Sydney Siege |      |      | Ferguson |      |      | Ottawa Shoot. |      |      | Germanwings |      |      | Putin Missing |      |      | Overall ASR |
|-------------------------|---------------|------|------|--------------|------|------|----------|------|------|---------------|------|------|-------------|------|------|---------------|------|------|-------------|
|                         | A             | C    | H    | A            | C    | H    | A        | C    | H    | A             | C    | H    | A           | C    | H    | A             | C    | H    |             |
| Proprietary Models      |               |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| GPT-4o-mini             | 81.7          | 0.83 | 67.4 | 92.1         | 0.82 | 70.3 | 79.5     | 0.81 | 59.0 | 86.7          | 0.85 | 78.2 | 74.5        | 0.85 | 60.0 | 66.7          | 0.82 | 53.3 | 83.1        |
| GPT-4o                  | 79.4          | 0.85 | 66.9 | 88.5         | 0.83 | 69.1 | 67.0     | 0.83 | 53.0 | 85.5          | 0.85 | 74.5 | 56.4        | 0.89 | 52.7 | 63.3          | 0.75 | 26.7 | 77.4        |
| GPT-4.1-nano            | 81.7          | 0.81 | 66.9 | 94.5         | 0.80 | 67.1 | 79.1     | 0.79 | 53.1 | 94.5          | 0.81 | 74.8 | 74.5        | 0.83 | 61.8 | 70.0          | 0.75 | 13.3 | 85.5        |
| GPT-4.1-mini            | 79.9          | 0.85 | 68.4 | 80.6         | 0.82 | 50.9 | 68.0     | 0.81 | 42.5 | 77.6          | 0.87 | 67.9 | 63.6        | 0.91 | 63.6 | 20.0          | 0.81 | 16.7 | 72.7        |
| GPT-4.1                 | 77.1          | 0.88 | 73.7 | 64.8         | 0.88 | 60.6 | 60.0     | 0.87 | 57.5 | 77.0          | 0.89 | 74.5 | 54.5        | 0.94 | 54.5 | 16.7          | 0.90 | 16.7 | 65.9        |
| Claude-3-Haiku          | 94.8          | 0.80 | 69.0 | 98.8         | 0.79 | 72.0 | 83.5     | 0.76 | 50.0 | 98.8          | 0.79 | 66.9 | 68.5        | 0.82 | 61.1 | 86.7          | 0.71 | 26.7 | 91.5        |
| Claude-3.5-Haiku        | 76.9          | 0.81 | 47.9 | 76.9         | 0.79 | 45.0 | 77.1     | 0.79 | 46.9 | 80.5          | 0.77 | 38.4 | 61.8        | 0.82 | 38.2 | 74.1          | 0.71 | 18.5 | 76.7        |
| Claude-4.5-Haiku        | 52.6          | 0.74 | 20.6 | 34.5         | 0.70 | 10.3 | 32.0     | 0.71 | 8.0  | 50.9          | 0.74 | 14.6 | 63.6        | 0.77 | 34.5 | 16.7          | 0.61 | 16.7 | 42.4        |
| Proprietary Avg.        | 78.0          | 0.82 | 60.1 | 78.8         | 0.80 | 55.7 | 68.3     | 0.80 | 46.2 | 81.4          | 0.82 | 61.2 | 64.7        | 0.85 | 53.3 | 51.8          | 0.76 | 23.6 | 74.4        |
| Open-Weights Models     |               |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| Qwen2.5-3B-Inst         | 54.7          | 0.88 | 51.2 | 59.4         | 0.88 | 53.1 | 52.3     | 0.88 | 49.2 | 64.4          | 0.88 | 60.0 | 45.3        | 0.90 | 43.4 | 31.0          | 0.78 | 20.7 | 55.4        |
| Qwen2.5-7B-Inst         | 67.8          | 0.86 | 62.1 | 82.4         | 0.84 | 68.5 | 62.0     | 0.85 | 56.5 | 64.0          | 0.86 | 54.3 | 65.5        | 0.86 | 61.8 | 36.7          | 0.80 | 26.7 | 67.1        |
| Qwen2.5-14B-Inst        | 71.4          | 0.85 | 59.5 | 81.3         | 0.82 | 62.6 | 60.2     | 0.83 | 46.9 | 85.9          | 0.85 | 69.3 | 53.9        | 0.89 | 50.0 | 69.0          | 0.75 | 27.6 | 71.9        |
| DS-R1-Distill-Qwen-1.5B | 71.0          | 0.76 | 50.8 | 81.0         | 0.76 | 55.6 | 61.1     | 0.70 | 37.6 | 72.5          | 0.74 | 47.5 | 79.1        | 0.73 | 51.2 | 75.0          | 0.62 | 33.3 | 71.6        |
| DS-R1-Distill-Qwen-7B   | 74.9          | 0.88 | 66.3 | 92.1         | 0.85 | 75.2 | 76.0     | 0.87 | 69.5 | 81.2          | 0.89 | 74.5 | 66.0        | 0.92 | 60.4 | 66.7          | 0.86 | 60.0 | 79.2        |
| DS-R1-Distill-Qwen-14B  | 77.0          | 0.86 | 64.9 | 83.6         | 0.85 | 71.5 | 74.9     | 0.84 | 64.3 | 88.5          | 0.88 | 84.2 | 54.5        | 0.91 | 52.7 | 36.7          | 0.74 | 20.0 | 76.8        |
| Open-Weights Avg.       | 69.3          | 0.85 | 58.9 | 80.4         | 0.83 | 65.2 | 64.6     | 0.83 | 54.0 | 76.0          | 0.85 | 64.6 | 61.4        | 0.86 | 53.3 | 54.5          | 0.76 | 33.6 | 70.6        |

# Effect of CoT prompting


- Advanced inference becomes an attack surface under adversarial cognitive manipulation

| Victim Model            | Hebdo |      |      | Sydney Siege |      |      | Ferguson |      |      | Ottawa Shoot. |      |      | Germanwings |      |      | Putin Missing |      |      | Overall ASR |
|-------------------------|-------|------|------|--------------|------|------|----------|------|------|---------------|------|------|-------------|------|------|---------------|------|------|-------------|
|                         | A     | C    | H    | A            | C    | H    | A        | C    | H    | A             | C    | H    | A           | C    | H    | A             | C    | H    |             |
| Proprietary Models      |       |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| GPT-4o-mini             | 81.7  | 0.83 | 67.4 | 92.1         | 0.82 | 70.3 | 79.5     | 0.81 | 59.0 | 86.7          | 0.85 | 78.2 | 74.5        | 0.85 | 60.0 | 66.7          | 0.82 | 53.3 | 83.1        |
| GPT-4o                  | 79.4  | 0.85 | 66.9 | 88.5         | 0.83 | 69.1 | 67.0     | 0.83 | 53.0 | 85.5          | 0.85 | 74.5 | 56.4        | 0.89 | 52.7 | 63.3          | 0.75 | 26.7 | 77.4        |
| GPT-4.1-nano            | 81.7  | 0.81 | 66.9 | 94.5         | 0.80 | 67.1 | 79.1     | 0.79 | 53.1 | 94.5          | 0.81 | 74.8 | 74.5        | 0.83 | 61.8 | 70.0          | 0.75 | 13.3 | 85.5        |
| GPT-4.1-mini            | 79.5  |      |      |              |      |      |          |      |      |               |      |      |             |      | 6.6  | 20.0          | 0.81 | 16.7 | 72.7        |
| GPT-4.1                 | 77.3  |      |      |              |      |      |          |      |      |               |      |      |             |      | 4.5  | 16.7          | 0.90 | 16.7 | 65.9        |
| Claude-3-Haiku          | 94.8  |      |      |              |      |      |          |      |      |               |      |      |             |      | 1.1  | 86.7          | 0.71 | 26.7 | 91.5        |
| Claude-3.5-Haiku        | 76.9  |      |      |              |      |      |          |      |      |               |      |      |             |      | 1.2  | 74.1          | 0.71 | 18.5 | 76.7        |
| Claude-4.5-Haiku        | 52.0  |      |      |              |      |      |          |      |      |               |      |      |             |      | 4.5  | 16.7          | 0.61 | 16.7 | 42.4        |
| Proprietary Avg.        | 78.4  |      |      |              |      |      |          |      |      |               |      |      |             |      | 6.3  | 51.8          | 0.76 | 23.6 | 74.4        |
| Open-Weights Models     |       |      |      |              |      |      |          |      |      |               |      |      |             |      |      |               |      |      |             |
| Qwen2.5-3B-Inst         | 54.7  |      |      |              |      |      |          |      |      |               |      |      |             |      | 1.4  | 31.0          | 0.78 | 20.7 | 55.4        |
| Qwen2.5-7B-Inst         | 67.8  | 0.86 | 62.1 | 82.4         | 0.84 | 68.5 | 62.0     | 0.85 | 56.5 | 64.0          | 0.86 | 54.3 | 65.5        | 0.86 | 61.8 | 36.7          | 0.80 | 26.7 | 67.1        |
| Qwen2.5-14B-Inst        | 71.4  | 0.85 | 59.5 | 81.3         | 0.82 | 62.6 | 60.2     | 0.83 | 46.9 | 85.9          | 0.85 | 69.3 | 53.9        | 0.89 | 50.0 | 69.0          | 0.75 | 27.6 | 71.9        |
| DS-R1-Distill-Qwen-1.5B | 71.0  | 0.76 | 50.8 | 81.0         | 0.76 | 55.6 | 61.1     | 0.70 | 37.6 | 72.5          | 0.74 | 47.5 | 79.1        | 0.73 | 51.2 | 75.0          | 0.62 | 33.3 | 71.6        |
| DS-R1-Distill-Qwen-7B   | 74.9  | 0.88 | 66.3 | 92.1         | 0.85 | 75.2 | 76.0     | 0.87 | 69.5 | 81.2          | 0.89 | 74.5 | 66.0        | 0.92 | 60.4 | 66.7          | 0.86 | 60.0 | 79.2        |
| DS-R1-Distill-Qwen-14B  | 77.0  | 0.86 | 64.9 | 83.6         | 0.85 | 71.5 | 74.9     | 0.84 | 64.3 | 88.5          | 0.88 | 84.2 | 54.5        | 0.91 | 52.7 | 36.7          | 0.74 | 20.0 | 76.8        |
| Open-Weights Avg.       | 69.3  | 0.85 | 58.9 | 80.4         | 0.83 | 65.2 | 64.6     | 0.83 | 54.0 | 76.0          | 0.85 | 64.6 | 61.4        | 0.86 | 53.3 | 54.5          | 0.76 | 33.6 | 70.6        |

# Downstream decision simulation

---

- Both Strategies remain highly vulnerable with DDR above 50%



# Conclusion

---

- Formalize "Cognitive Colusion Attack"
- Propose "Generative Montage" framework
- Introduce "CoPHEME" dataset and conduct extensive experiments

# Opinion

---

- 장점
  - 새로운 공격 관점 제시
  - 명확한 프레임워크
- 아쉬운점
  - 사람을 대상으로 한 분석 내용 부재